

# AEGIRA Technical Build Playbook

# AEGIRA Technical Build Playbook

*(Time-aware longitudinal workplace risk infrastructure — build-from-scratch guide for an engineer)*

---

## 0) North Star: what we're building

AEGIRA is an enterprise system that continuously captures **worker check-ins + context (shift/roster/leave)** and turns them into **time-aware longitudinal signals** to detect emerging risk early. It must be:

- **Time-correct** (event-time first, not “when we received it”)
  - **Longitudinal** (per-person baselines, trends, deviations)
  - **Auditable** (no silent overwrites; every flag explainable)
  - **Enterprise-ready** (RBAC, privacy, multi-tenant, scale)
- 

## 1) Users and workflows

### Users

- **Worker**: completes check-ins; can see own status, reminders, and “support options.”
- **Supervisor / Leading Hand**: sees *actionable* flags (not raw psych data) for their team.
- **HSE / WHS**: sees site-level risk, trends, episodes, outcomes.
- **RTW / Rehab / Case Manager**: manages risk episodes/cases, actions, follow-ups.
- **Org Admin**: config, roster integration, forms, permissions, retention/deletion.

### Core workflow (happy path)

1. Worker is assigned to site/team and has a shift schedule.
2. System generates a **CheckInPrompt** for a shift window (pre-start / mid / end).
3. Worker submits **CheckInResponse** (event\_time, answers, optional notes).
4. System computes **features** vs personal baseline (deviation, trend, missingness).
5. Risk engine creates/uploads a **RiskEpisode** and may open a **Case**.
6. Case gets routed to the right queue (Supervisor / HSE / RTW) with recommended actions.

7. Actions/outcomes logged; model/rule versions recorded; reporting updates.

## Edge-case workflow (critical)

- Worker **misses** a check-in → system records a “missed” event at the end of the window.
  - Worker submits **late** → system records the late response as its own event and **reconciles** the missed event (without deleting history).
  - Worker is **off-shift / on leave** → no missed event should be generated; or if already generated, it must be corrected with an auditable “status override” event.
- 

## 2) Architecture: components and responsibilities

### Recommended architecture (pragmatic, production-ready)

#### Frontend

- Worker app (mobile/web)
- Admin + supervisor dashboard (web)

#### API & Auth

- API Gateway
- Auth (OIDC/SAML for enterprise, passwordless for workers if needed)
- RBAC enforcement at API layer

#### Event ingestion (event-time first)

- Events API (append-only)
- Event bus/stream (or queue) for async processing
- Idempotency + dedupe

#### Storage

- **OLTP DB** (entities, prompts, cases, permissions) – Postgres
- **Event store** (append-only log of check-ins/status changes) – Postgres (append-only table) or Kafka + compacted storage

- **Analytics warehouse** (dashboards, cohort reporting) – BigQuery/Snowflake/Redshift (optional early)

## Compute

- Feature engine (batch + streaming)
- Baseline engine (per-person)
- Risk scoring engine (rules now, ML-ready later)

## Case management

- Episode builder
- Routing + notifications (email/SMS/Teams/Slack)
- Action/outcome logging

## Observability + audit

- Central logging + traces + metrics
- Audit export pipeline (raw + derived + explanation)

## Text diagram

Worker App → API Gateway/Auth → Events API → Event Store  
 Event Store → Stream Processor → Feature/Baseline Store → Risk Engine → Episodes/Cases  
 DB → Dashboards/Notifications  
 OLTP DB provides Workers/Teams/Rosters/Forms/Permissions to all services.

---

## 3) Data principles (this is what makes it “time-aware”)

### Two timestamps everywhere

- `event_time`: when it *happened* (worker submitted, shift started/ended, leave applied)
- `ingested_at`: when AEGIRA *received* it

All analytics, baselines, and missed-check logic use `event_time`.

### Immutable event log + “state projections”

- Event log is append-only (never overwrite reality).

- Current “truth” is derived via **projections**:
  - “Latest check-in for prompt X”
  - “Worker availability status on date/time”
  - “Current risk episode state”

## Idempotency

Every event from a client/integration must include:

- `event_id` (UUID) + `source` + `source_event_id` (optional)
  - A dedupe key: (`tenant_id, event_id`) or (`tenant_id, source, source_event_id`)
- 

## 4) Core domain model (entities)

### Org structure

- **Tenant**: enterprise customer
- **Site**
- **Team**
- **Worker**: belongs to tenant; mapped to site/team; has privacy attributes
- **Roster/Shift**: shift windows in a timezone, source = manual or integration
- **Leave/OffShift**: explicit unavailability

### Check-in system

- **Form**: versioned questions (fatigue, stress, pain, sleep, confidence, etc.)
- **CheckInPrompt**: a scheduled request tied to a worker and a shift window
- **CheckInResponse**: a submission tied to a prompt (or “ad hoc”)

### Risk & workflow

- **RiskSignal**: computed result for a specific response/window (deviation metrics + explainability)
- **RiskEpisode**: a time-bounded cluster of risk signals (per worker)
- **Case**: human workflow object (owner, status, actions, outcomes)

## Audit and explainability

- **DecisionLog**: records risk engine decisions, rule/model version, feature snapshot hashes
- 

## 5) Event schemas (append-only) — the “source of truth”

You will store every event in a single table/stream, then project.

### Base event envelope (every event uses this)

```
{  
  "event_id": "uuid",  
  "tenant_id": "uuid",  
  "event_type": "checkin.response_submitted",  
  "event_time": "2026-02-12T07:12:03+11:00",  
  "ingested_at": "2026-02-12T07:12:05+11:00",  
  "actor": { "type": "worker", "id": "worker_uuid" },  
  "source": "worker_app",  
  "source_event_id": "optional",  
  "schema_version": 1,  
  "payload": {}  
}
```

### Key event types (minimum viable)

#### 1) roster.shift\_upserted

```
{  
  "shift_id": "uuid",  
  "worker_id": "uuid",  
  "site_id": "uuid",  
  "timezone": "Australia/Perth",  
  "start_time": "2026-02-12T06:00:00+08:00",  
  "end_time": "2026-02-12T18:00:00+08:00",  
  "source": "integration|manual"  
}
```

## **2) worker.availability\_set**

*(off-shift/leave/sick)*

```
{  
  "worker_id": "uuid",  
  "status": "on_leave|sick|off_shift|available",  
  "effective_from": "2026-02-12T00:00:00+11:00",  
  "effective_to": "2026-02-14T23:59:59+11:00",  
  "reason": "Annual leave",  
  "source": "manual|integration"  
}
```

## **3) checkin.prompt\_scheduled**

```
{  
  "prompt_id": "uuid",  
  "worker_id": "uuid",  
  "shift_id": "uuid",  
  "form_id": "uuid",  
  "window": {  
    "opens_at": "2026-02-12T05:15:00+08:00",  
    "closes_at": "2026-02-12T06:05:00+08:00"  
  },  
  "prompt_kind": "pre_shift|mid_shift|post_shift"  
}
```

## **4) checkin.response\_submitted**

```
{  
  "response_id": "uuid",  
  "prompt_id": "uuid",  
  "worker_id": "uuid",  
  "submitted_at": "2026-02-12T06:12:03+08:00",  
  "answers": {  
    "fatigue": 7,  
    "stress": 5,  
    "pain": 3,  
    "sleep_quality": 4,  
    "readiness": 5  
  },  
}
```

```
    "note": "Shoulder a bit tight",
    "client_metadata": { "app_version": "1.0.9" }
}
```

## 5) **checkin.window\_missed**

*(system-generated at window close)*

```
{
  "missed_id": "uuid",
  "prompt_id": "uuid",
  "worker_id": "uuid",
  "window_closed_at": "2026-02-12T06:05:00+08:00",
  "reason": "no_submission",
  "availability_snapshot": "available|off_shift|on_leave"
}
```

## 6) **risk.signal\_computed**

```
{
  "signal_id": "uuid",
  "worker_id": "uuid",
  "response_id": "uuid",
  "computed_for_time": "2026-02-12T06:12:03+08:00",
  "features": {
    "fatigue_z": 2.1,
    "stress_z": 1.2,
    "pain_delta": 2,
    "missingness_7d": 0.33,
    "volatility_14d": 1.7
  },
  "explanations": [
    { "feature": "fatigue_z", "why": "fatigue 7 vs baseline 3.8" },
    { "feature": "missingness_7d", "why": "2 missed of last 6 prompts" }
  ],
  "engine_version": "ruleset_v3.2"
}
```

## 7) **risk.episode\_updated / case.opened / case.action\_logged**

Keep episodes and case changes as events too (optional in v1; at least log them in DB + DecisionLog).

---

## 6) The “missed vs late vs overwrite” rule (exactly how to implement)

### Golden rule

**Never overwrite raw events.** Use reconciliation via projections.

#### What happens when a check-in is missed?

- At `window.closes_at`, the system emits `checkin.window_missed` (append-only).

#### What happens when a worker submits late?

- The worker app submits `checkin.response_submitted` with `submitted_at` (`event_time`).
- The projection layer marks:
  - Prompt status = `completed_late`
  - Links response to prompt
  - Keeps the missed event in history but sets prompt’s *current state* to “`completed_late`”

#### If a worker was actually off-shift

- An availability event `worker.availability_set` arrives (or shift is removed).
- Projection re-evaluates prompts in that time window:
  - Prompt current state becomes `excused`
  - Any existing `window_missed` remains in event log but is now “explained/excused” in the current state.

**Result:** You retain a complete audit trail while the “truth” used for metrics is correct.

---

# 7) Storage design (tables you actually build)

## A) OLTP tables (Postgres)

- `tenants, sites, teams, workers`
- `forms, form_versions, questions`
- `shifts` (optional if not purely event-derived)
- `prompts` (current state + references)
- `cases, case_actions, case_outcomes`
- `rbac_roles, rbac_permissions, user_accounts`

## B) Event store table (append-only)

`events`

- `event_id` (pk)
- `tenant_id`
- `event_type`
- `event_time` (`timestamptz`)
- `ingested_at` (`timestamptz`)
- `actor_type, actor_id`
- `source, source_event_id`
- `payload` (`jsonb`)
- `schema_version`
- Indexes:
  - `(tenant_id, event_time)`
  - `(tenant_id, actor_id, event_time)`
  - `(tenant_id, event_type, event_time)`
  - Dedupe: `(tenant_id, source, source_event_id)` unique when present

## C) Projections (materialized views or tables updated by processors)

- `prompt_state`
  - `prompt_id, worker_id, shift_id, opens_at, closes_at`

- `state`:  
    scheduled|completed\_on\_time|completed\_late|missed|excused
  - `response_id` (nullable)
  - `last_updated_at`
  - `worker_daily_state`
    - worker availability timeline snapshots (or query from events efficiently)
  - `feature_store`
    - worker\_id, computed\_at, feature\_name, value, window, version
  - `risk_episode_state`
    - worker\_id, episode\_id, status, severity, opened\_at, last\_signal\_at
- 

## 8) API surface (minimum endpoints)

### Worker

- `GET /worker/me`
- `GET /worker/prompts?from=&to=`
- `POST /worker/checkins` (*submits response; idempotent*)
- `GET /worker/history` (*optional*)

### Admin / Ops

- `POST /admin/forms`
- `POST /admin/shifts:upsert` (*or integration webhook*)
- `POST /admin/availability`
- `GET /admin/workers`
- `GET /admin/prompts`
- `GET /admin/cases`
- `POST /admin/cases/{id}/actions`

### System/Integration

- `POST /events` (*generic ingestion endpoint; recommended internally even if you also expose “friendly” endpoints*)
- `GET /exports/raw-events?worker_id=&from=&to=` (*audit export*)

All POSTs **idempotent** via header `Idempotency-Key`.

---

## 9) Time-aware computation (baseline + features)

### Baseline strategy (v1: robust + explainable)

For each metric (fatigue, stress, pain, sleep, readiness):

- Use a **rolling window** (e.g., last 28 days of valid responses)
- Baseline = **median** (robust to outliers)
- Dispersion = **MAD** (median absolute deviation) or IQR
- Convert current value to deviation score:
  - `z_robust = (x - median) / (1.4826 * MAD + epsilon)`

### Features to compute (minimum)

#### Level & deviation

- `metric_z_robust` per dimension

#### Trend

- 7d slope (linear regression on last N points)
- 14d slope

#### Volatility

- rolling standard deviation / MAD over 14d

#### Missingness / compliance

- missed rate over 7d / 28d
- streak length of misses

#### Context-aware features

- “post-nightshift penalty” (optional)
- shift length, consecutive shifts, overtime flags (if available)

## Under-reporters (“stoic workers”)

Use **change from baseline** and **within-person volatility**, not absolute thresholds.

- A worker always reporting “2/10” can still show risk if:
    - volatility rises
    - misses increase
    - subtle upward drift emerges
- 

## 10) Risk engine (rules first, ML-ready)

### Output objects

- `RiskSignal` (per response/window)
- `RiskEpisode` (cluster of signals)
- Optional `Case` (human workflow)

### Ruleset v1 (example)

Trigger a **RiskSignal severity**:

- **High**
  - Any `fatigue_z > 2.5 OR stress_z > 2.5 OR pain_delta >= 3`
  - OR `2+ dimensions with z > 2.0`
- **Medium**
  - `fatigue_z > 1.8` or trend slope rising + volatility rising
- **Compliance-risk**
  - `missed_rate_7d > 0.4 AND worker availability = available`

### Episode logic

- Open episode if:
  - High signal, or
  - 2 medium within 7 days, or
  - Compliance-risk sustained 2+ windows
- Close episode if:
  - No medium/high signals for 14 days AND compliance normalizes, OR
  - Case resolved + outcome logged

## **Routing logic (enterprise-friendly)**

- Supervisor sees: “Worker needs a check-in / support”
  - HSE/RTW sees: reason codes + suggested next step
  - Limit raw metric exposure based on role
- 

## **11) Dashboards and metrics (what workplace teams need)**

### **For supervisors**

- Team today: who is flagged, who missed, who is excused
- “Next best action” suggestions
- Compliance view (not shaming; operational)

### **For HSE/RTW**

- Active episodes by site/team
- Time-to-intervention
- Episode outcomes (resolved, escalated, injury prevented proxy)

### **For exec**

- Leading indicators: risk trend, compliance trend
  - Lagging indicators: incidents/claims integration later
- 

## **12) Data trust, exportability, auditability (non-negotiable)**

### **Every flag must be explainable**

Store, for each `risk.signal_computed`:

- Features used + values
- Baseline snapshot identifiers
- Rule/model version
- Human-readable reasons (“fatigue 7 vs baseline 3.8”)

## Raw export

Build a deterministic export that includes:

- Raw events (checkins, missed, availability, roster changes)
- Projections (prompt\_state)
- Computed signals (features + reasons)
- Episode/case timeline

## Reproducibility

Version everything:

- Form versions
- Ruleset versions
- Feature computation versions

If you re-run history, you must be able to reproduce *the same decision* given the same version inputs.

---

## 13) Security + privacy (enterprise baseline)

### RBAC roles (minimum)

- Worker (self only)
- Supervisor (team-level, redacted)
- HSE/WHS (broader, controlled)
- Case Manager (cases + signals)
- Org Admin (config)
- Platform Admin (internal, restricted)

### Privacy patterns

- Separate “identifiers” (name, phone) from health-ish responses
- Role-based redaction: supervisors don’t see raw scores unless allowed
- Tenant isolation in every query

## Retention and deletion

- Retention policies per tenant
  - Delete worker: tombstone identity + remove direct identifiers; keep aggregated stats if permitted
- 

# 14) Step-by-step build plan (milestones + acceptance criteria)

## Phase 1 — Foundations: events + identity + RBAC

### Build

- Tenant/org model, workers, teams, sites
- Auth + RBAC
- Append-only `events` table + `/events` ingestion endpoint
- Idempotency + dedupe

### Acceptance

- You can ingest 1M events without duplicates
  - You can retrieve worker event history efficiently by time range
- 

## Phase 2 — Shift-aware prompting + missed logic

### Build

- Shift ingestion (`roster.shift_upserted`)
- Prompt scheduler (creates `checkin.prompt_scheduled`)
- Window close job emits `checkin.window_missed`
- Availability events + excusal projection

### **Acceptance**

- Prompt\_state correctly shows: scheduled / missed / completed\_on\_time / completed\_late / excused
  - Late submission never deletes missed history
- 

## **Phase 3 — Baselines + feature engine**

### **Build**

- Baseline computation per worker per metric (rolling median + MAD)
- Feature computation pipeline on new responses and missed events
- Feature store tables

### **Acceptance**

- For any response, you can show baseline and deviation in an export
  - Backfill works deterministically
- 

## **Phase 4 — Risk engine + episodes + case workflow**

### **Build**

- Ruleset engine
- RiskSignal creation event
- Episode state machine
- Case creation/routing + action logging

### **Acceptance**

- Given test histories, engine opens/closes episodes correctly
  - Every case has reasons + feature snapshot references
- 

## **Phase 5 — Dashboards + reporting + exports**

### **Build**

- Supervisor dashboard views

- HSE/RTW dashboard views
- Export endpoints (raw + derived)
- Metrics: compliance, risk trend, episode counts

#### Acceptance

- Stakeholder can validate flags from raw export alone
  - Basic cohort reporting works per site/team/time window
- 

## Phase 6 — Enterprise hardening

#### Build

- SSO/SAML, audit logs, encryption, secrets management
- Monitoring, alerting, rate limits
- Multi-tenant performance testing
- Data retention/deletion jobs

#### Acceptance

- Pass security review checklist
  - Meet agreed SLOs; can scale per tenant
- 

## 15) Testing strategy (what to test so you don't get destroyed in pilots)

### Unit tests

- event-time ordering
- dedupe/idempotency
- baseline math (median/MAD)
- feature correctness
- ruleset triggers

### Scenario tests (must-have)

1. Missed then late submission → final state `completed_late`, history includes `missed`

2. Shift removed after missed → final state **excused**, history preserved
  3. Availability set to leave → prompts not generated, or excused if already created
  4. Multi-timezone (WA vs NSW) → windows correct
  5. Backfill re-run with same versions → identical outcomes
- 

## 16) “Exact process” for an engineer starting tomorrow

1. Implement **events ingestion** + append-only store + idempotency.
  2. Implement **rosters + availability** events and projections.
  3. Implement **prompt scheduling** and **missed window** generation.
  4. Implement **check-in submission** and “late vs on-time” resolution in projections.
  5. Implement **baseline engine** (rolling median/MAD) + feature pipeline.
  6. Implement **risk signals** (rules v1) + decision logging.
  7. Implement **episodes + cases** with routing and action/outcome tracking.
  8. Implement **exports** that include raw + derived + explanations.
  9. Build dashboards with strict RBAC redaction.
  10. Harden for enterprise (SSO, audit logs, retention, deletion, scaling).
-

# What we have/what we need

Good — this is exactly the right moment to audit the structure before it hardens.

You currently have **1 row, 20 columns**, and structurally this is a **flat check-in snapshot table**.

That is **not yet a longitudinal, time-aware risk system**.

Below is a CTO-level audit against the AEGIRA Technical Build Playbook.

---



## What You Currently Have

Columns:

```
id  
company_id  
person_id  
event_id  
check_in_date  
hours_slept  
sleep_quality  
stress_level  
physical_condition  
notes  
readiness_score  
readiness_level  
sleep_score  
stress_score  
physical_score  
created_at  
pain_level  
pain_location  
physical_condition_notes  
pain_score
```

This is essentially:

“One row = one check-in submission”

That's fine for v0.

It is **not enough** for AEGIRA's longitudinal risk infrastructure.

---

# Critical Gaps (Against AEGIRA Architecture)

## 1 No True Event-Time System

You have:

- `check_in_date`
- `created_at`

You do NOT have:

- `event_time` (true submission time in timezone)
- `ingested_at`
- `timezone`
- shift reference
- window open/close
- late vs on-time logic

Right now:

- You cannot determine if this was late.
- You cannot determine if a check-in was missed.
- You cannot detect off-shift excusal.
- You cannot handle multi-timezone enterprises.

This is a major architectural gap.

---

## 2 No Shift / Prompt Context

There is:

- No `shift_id`
- No `prompt_id`
- No `window.opens_at`
- No `window.closes_at`

- No `prompt_kind` (pre/mid/post shift)

Without that:

You cannot:

- Generate missed events
- Detect non-compliance
- Reconcile late submissions
- Exclude off-shift workers
- Perform time-aware compliance analytics

Right now this is just a survey table.

---

## 3 No Missed Check-In Events

There is no structure for:

- `checkin.window_missed`
- availability snapshot
- excused status
- reconciliation state

This means:

If a worker misses a check-in, nothing is recorded.

That is catastrophic for:

- Compliance trend detection
- Behavioural drift detection
- Stoic under-reporters
- Risk emergence via disengagement

AEGIRA MUST log missed events.

---

## 4 No Event-Sourcing

This table is state-based, not event-based.

You have:

- `event_id` (but unclear what that represents)
- `id` (primary key)

But:

There is no append-only event store.  
No schema version.  
No actor metadata.  
No source.  
No immutability guarantees.

If someone edits a check-in row, you lose history.

AEGIRA requires:

Append-only event log + projections.

---

## 5 No Longitudinal Computation Layer

This file stores:

Raw responses + derived scores.

It does NOT store:

- Baseline snapshot
- Rolling median
- MAD
- z-score deviation
- volatility
- trend slope
- missingness rate
- compliance rate
- feature version
- ruleset version

Which means:

You cannot explain why someone is flagged.

And that kills enterprise trust.

---

## 6 No Availability / Off-Shift Tracking

There is no:

- leave event
- sick event
- off-shift status
- effective\_from / effective\_to

So if a worker is on leave, your system would:

Incorrectly mark them as non-compliant.

This must be event-driven.

---

## 7 No Episode or Case Model

No:

- risk\_signal
- risk\_episode
- case
- case\_action
- outcome

So currently:

This is not a workplace risk system.

It is a check-in storage table.

---



## What You Did Right

To be fair:

- `person_id`
- `company_id`
- numeric domain fields
- separation of raw inputs and computed scores
- `created_at`

This is a solid starting capture schema.

But it is not yet longitudinal infrastructure.

---



## What Must Change (Exact Structural Upgrades)

### 1 Convert to Event-Based Capture

Instead of storing only rows in `check_ins_rows`, you need:

#### A) Append-only `events` table

Add:

```
event_id (uuid)
tenant_id
event_type
event_time (timestamptz)
ingested_at (timestamptz)
actor_type
actor_id
source
schema_version
payload (jsonb)
```

Your current row becomes:

```
event_type = "checkin.response_submitted"
```

Everything else moves into payload.

---

### 2 Add Prompt + Shift Awareness

You must introduce:

```
shift_id  
prompt_id  
window.opens_at  
window.closes_at  
prompt_kind  
timezone
```

Without this you cannot detect missed vs late.

---

## 3 Add Missed Event Generation

At window\_closes\_at:

If no response exists:

Emit event:

checkin.window\_missed

This must be stored as a row.

Never implied. Always recorded.

---

## 4 Separate Raw Inputs from Derived Scores

Currently you store:

- sleep\_score
- stress\_score
- readiness\_score
- physical\_score

Question:

Are these raw user inputs?  
Or derived?

If derived, they must be:

- Versioned
- Recomputable
- Logged with feature version

Otherwise historical flags cannot be reproduced.

---

## 5 Add Feature Store

You need a table:

worker\_id  
computed\_at  
feature\_name  
feature\_value  
window  
version

Example:

fatigue\_z = 2.1  
volatility\_14d = 1.7  
missing\_rate\_7d = 0.33

Without this, you are not longitudinal.

---

## 6 Add Availability Model

Introduce event:

worker.availability\_set  
effective\_from  
effective\_to  
status

This is mandatory before pilots.

---



# Most Important Strategic Point

Right now your data structure supports:

“How did someone feel today?”

AEGIRA must answer:

“Is this worker deviating from their personal baseline over time in a way that predicts risk?”

That requires:

- Event-time correctness
- Shift context
- Missed detection
- Baseline modelling
- Deviation tracking
- Versioned risk decisions
- Audit exports

You currently have none of those.

That's not a criticism — that's just architectural reality.

---



## Immediate Upgrade Roadmap (Practical)

If I were your CTO today, I would say:

Step 1 (this week):

- Introduce event store.
- Convert check-in submissions into events.
- Add event\_time + timezone.

Step 2:

- Add shift ingestion + prompt scheduling.
- Generate missed events.

Step 3:

- Add baseline computation engine.
- Store deviation metrics separately.

Step 4:

- Implement risk signals + episodes.

Only after that are you building AEGIRA.

---

## Final CTO Verdict

Your CSV is:

- A check-in capture table
- Not longitudinal
- Not time-aware
- Not risk infrastructure
- Not enterprise-grade
- Not audit-ready

But it is a clean starting foundation.

production-ready redesign

Here's a **production-ready redesign** of your current `check_ins_rows.csv` into a **time-aware, longitudinal, auditable event log CSV** (what AEGIRA actually needs).

You'll end up with **one primary export**: `aegira_event_log.csv`.

(You *can* add supporting CSVs later, but this one is the non-negotiable foundation.)

---

## 1) New CSV: `aegira_event_log.csv` (single source of truth)

### Why this format

- **Append-only**: no overwrites
- Handles **late check-ins, missed windows, off-shift/leave, and roster changes**
- Keeps both **event\_time** and **ingested\_at**
- Supports **reproducibility** (schema\_version, source, idempotency keys)

### Columns (header)

These columns are designed so multiple event types can live in one CSV (unused fields are blank).

```
event_id,tenant_id,worker_id,event_type,event_time,ingested_at,timezone,source,source_event_id,schema_version,  
prompt_id,shift_id,prompt_kind,window.opens_at,window.closes_at,submitted_at,is_late,late_by_seconds,  
availability_status,availability_effective_from,availability_effective_to,availability_reason,  
hours_slept,sleep_quality,stress_level,physical_condition,pain_level,pain_location,readiness_level,  
notes,physical_condition_notes,  
sleep_score,stress_score,physical_score,pain_score,readiness_score,  
dedupe_key,correlation_id
```

### What each event\_type means

Minimum set you should emit:

- `checkin.response_submitted` (worker submitted a check-in)
- `checkin.window_missed` (system-generated at window close if no submission)
- `worker.availability_set` (leave/off-shift/sick/available)

- `roster.shift_upserted` (shift added/updated; optional early but strongly recommended)
- 

## 2) Exact rules for missed vs late vs off-shift (how this CSV behaves)

### Missed check-in

At `window_closes_at`, if no response exists:

- Write **one new row**: `event_type=checkin.window_missed`
- Do **not** “assume” missed in reporting — it must be an actual recorded event

### Late check-in

If worker submits after window close:

- Write **one new row**: `event_type=checkin.response_submitted`
- Set `is_late=true`, fill `late_by_seconds`
- The earlier `checkin.window_missed` row stays forever (audit trail)
- Your projections later decide the prompt’s current state = `completed_late`

### Off-shift / leave

When a worker is not expected to check in:

- Write **availability** event(s): `event_type=worker.availability_set`
  - Then missed events can be **excused** in projections (but still remain if they were already emitted)
- 

## 3) Mapping from your current CSV to the new one

Your current columns map like this:

- `company_id` → `tenant_id`
- `person_id` → `worker_id`
- `event_id` → `correlation_id` (*or source\_event\_id if it came from client*)

- `check_in_date` → NOT enough on its own; becomes `event_time` only if you also store a time + timezone
  - `created_at` → `ingested_at` (or keep as `ingested_at` if it truly is “received time”)
  - all check-in answers/scores → same names in new CSV

## **Important change:**

You must start capturing a true submission timestamp:

- `submitted_at` (event\_time for the response)
  - plus `timezone`

#### **4) Example rows (what your export will look like)**

### **A) Worker submitted on time**

event\_id,tenant\_id,worker\_id,event\_type,event\_time,ingested\_at,timezone,source,source\_event\_id,schema\_version,prompt\_id,shift\_id,prompt\_kind>window.opens\_at>window.closes\_at,submitted\_at,is\_late,late\_by\_seconds,availability\_status,availability\_effective\_from,availability\_effective\_to,availability\_reason,hours\_slept,sleep\_quality,stress\_level,physical\_condition,pain\_level,pain\_location,readiness\_level,notes,physical\_condition\_notes,sleep\_score,stress\_score,physical\_score,pain\_score,readiness\_score,dedupe\_key,correlation\_id  
66ba70e9-698a-49e6-812d-058f555dec1a,f0ee43c6-5b6e-4946-801e-0bb15eef6d26,1d2ef946-2e50-481a-b33f-7d3c4cdf438b,checkin.response\_submitted,2026-02-12T14:02:00+11:00,2026-02-12T14:02:00+11:00,Australia/Sydney,worker\_app,,1,pr\_001,sh\_001,pre\_shift,2026-02-12T13:50:00+11:00,2026-02-12T14:10:00+11:00,2026-02-12T14:02:00+11:00,false,,available,,,7,7,5,6,0,,GREEN,,"85","50","60","100","74",tenant:f0ee43c6|worker:1d2ef946|prompt:pr\_001|submitted:2026-02-12T14:02:00+11:00,66ba70e9-698a-49e6-812d-058f555dec1a

### **B) System recorded a missed window**

### C) Worker marked as on leave (excuses prompts)

```
event_id,tenant_id,worker_id,event_type,event_time,ingested_at,timezone,source,source_event_id,schema_version,prompt_id,shift_id,prompt_kind,window.opens_at,window.closes_at,submitted_at,is_late,late_by_seconds,availability_status,availability_effective_from,availability_effective_to,availability_reason,hours_slept,sleep_quality,stress_level,physical_condition,pain_level,pain_location,readiness_level,notes,physical_condition_notes,sleep_score,stress_score,physical_score,pain_score,readiness_score,dedupe_key,correlation_id
avail_1001,f0ee43c6-5b6e-4946-801e-0bb15eef6d26,1d2ef946-2e50-481a-b33f-7d3c4cdf438b,worker.availability_set,2026-02-12T00:00:00+11:00,2026-02-12T08:01:00+11:00,Australia/Sydney,admin_dashboard,,1,,,,,,on_leave,2026-02-12T00:00:00+11:00,2026-02-14T23:59:59+11:00,Annual leave,,,,,,,,,,tenant:f0ee43c6|worker:1d2ef946|avail:2026-02-12,avail_1001
```

---

## 5) If you want the “proper” enterprise version (recommended)

You’ll eventually split into **supporting CSVs** for performance and clarity:

- `workers.csv` (worker, team, site)
- `shifts.csv` (shift windows)
- `prompts.csv` (prompt schedule + current state)
- `event_log.csv` (append-only truth)
- `feature_store.csv` (computed longitudinal features)
- `risk_signals.csv` (flags + explanations)
- `cases.csv + case_actions.csv`

But you can start with `event_log.csv` immediately and still be correct.

---

## 6) The one change your developer must make today

Right now you only have `check_in_date` and `created_at`.

To become time-aware, you must capture:

- `submitted_at` (with timezone)
- `window.opens_at / window.closes_at`

- `prompt_id` + `prompt_kind`
- and you must **emit missed rows**

Without those, “longitudinal risk” will always be guesswork.

---

If you want, I can also generate an updated CSV file (with headers + example rows) you can hand to your engineer as the target export spec—just say “generate the file”.

# Check In.csv Example

 Check In.csv